

Name:P.Asha Belcilda

Roll no:225229104

## Lab5: Diabetes Classification using Logistic Regression

### Step1.[Understand Data]

```
In [3]: import pandas as pd
```

```
In [4]: data=pd.read_csv("diabetes.csv")
print(data)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	
10	4	110	92	0	0	37.6	
11	10	168	74	0	0	38.0	
12	10	139	80	0	0	27.1	
13	1	189	60	23	846	30.1	
14	5	166	72	19	175	25.8	
15	7	100	0	0	0	30.0	
16	0	118	84	47	230	45.8	
17	7	107	74	0	0	29.6	

```
In [5]: #head
data.head()
```

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6	0.627	
1	1	85	66	29	0	26.6	0.351	
2	8	183	64	0	0	23.3	0.672	
3	1	89	66	23	94	28.1	0.167	
4	0	137	40	35	168	43.1	2.288	

```
In [7]: #shape
data.shape
```

```
Out[7]: (768, 9)
```

```
In [8]: #columns
data.shape[1]
```

```
Out[8]: 9
```

```
In [11]: #dtype
data.dtypes
```

```
Out[11]: Pregnancies      int64
Glucose      int64
BloodPressure  int64
SkinThickness int64
Insulin       int64
BMI           float64
DiabetesPedigreeFunction float64
Age           int64
Outcome       int64
dtype: object
```

```
In [12]: #info
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose          768 non-null int64
BloodPressure    768 non-null int64
SkinThickness    768 non-null int64
Insulin          768 non-null int64
BMI              768 non-null float64
DiabetesPedigreeFunction 768 non-null float64
Age              768 non-null int64
Outcome          768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [52]: #value_counts.  
data.count()
```

Out[52]: Pregnancies 768  
Glucose 768  
BloodPressure 768  
SkinThickness 768  
Insulin 768  
BMI 768  
DiabetesPedigreeFunction 768  
Age 768  
Outcome 768  
dtype: int64

Step2.[Build Logistic Regression Model]

```
In [32]: X=data.drop('Outcome',axis=1)
```

```
In [34]: X.head()
```

Out[34]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6		0.627
1	1	85	66	29	0	26.6		0.351
2	8	183	64	0	0	23.3		0.672
3	1	89	66	23	94	28.1		0.167
4	0	137	40	35	168	43.1		2.288

```
In [37]: y=data['Outcome'].values
y
```

```
Out[37]: array([1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0,
1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1,
1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1,
1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0,
1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0,
0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0,
1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0,
0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0,
0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1,
0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0,
0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1,
0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0,
0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0,
1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0],
dtype=int64)
```

```
In [40]: from sklearn.model_selection import StratifiedShuffleSplit
a=StratifiedShuffleSplit(n_splits=4,test_size=0.25,random_state=42)
```

```
In [41]: a.get_n_splits(X,y)
```

```
Out[41]: 4
```

```
In [47]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,stratify=y,test_size=.25,random_state=42)
```

```
In [48]: from sklearn.linear_model import LogisticRegression
```

```
In [53]: LOR=LogisticRegression(penalty='l2',C=10.0)
LOR=LOR.fit(X_train,y_train)
```

```
In [54]: y_pred=LOR.predict(X_test)
y_pred
```

```
Out[54]: array([0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0,
        0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
        1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1,
        0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
        0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0], dtype=int64)
```

### Step-3. [Predict on a new sample]

```
In [58]: new=LOR.predict([[6,200,90,10,25,23.3,.672,42]])
if new==0:
    print("Non-diabetic patient",new)
else:
    print("Diabetic patient",new)
```

Diabetic patient [1]

### Step3.[Compute Classification Metrics]

```
In [60]: #accuracy
def accuracy(actual,pred):
    return sum(actual==pred)/float(actual.shape[0])
```

```
In [61]: accuracy_score=y_test,y_pred)
accuracy_score
```

```
Out[61]: 0.734375
```

```
In [62]: #precision
from sklearn.metrics import precision_score
print(precision_score(y_test,y_pred))
```

0.6481481481481481

```
In [63]: #recall
from sklearn.metrics import recall_score
print(recall_score(y_test,y_pred))
```

0.5223880597014925

```
In [68]: #AUC scores
from sklearn.metrics import roc_auc_score
print(roc_auc_score(y_test,y_pred))
```

0.6851940298507463


## Step-4. [Understand Correlation]

```
In [69]: from sklearn.metrics import confusion_matrix
b=confusion_matrix(y_test,y_pred)
b
```

```
Out[69]: array([[106, 19],
               [ 32, 35]], dtype=int64)
```

```
In [70]: import seaborn as sns
sns.heatmap(b, annot=True)
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x2434875f5c0>
```



## Step-5. [Normalization using MinMaxScaler and rebuild LOR]

```
In [75]: from sklearn.preprocessing import MinMaxScaler
mm=MinMaxScaler()
mm_X_train=mm.fit_transform(X_train)
mm_X_train
```

```
Out[75]: array([[0.05882353, 0.6080402 , 0.63934426, ..., 0.58122206, 0.07884187,
                0.11666667],
               [0.70588235, 0.44221106, 0.60655738, ..., 0.52608048, 0.13095768,
                0.45          ],
               [0.05882353, 0.54271357, 0.49180328, ..., 0.5290611 , 0.14743875,
                0.05          ],
               ...,
               [0.05882353, 0.48743719, 0.57377049, ..., 0.56780924, 0.0596882 ,
                0.15          ],
               [0.52941176, 0.7839196 , 0.70491803, ..., 0.51117735, 0.4922049 ,
                0.35          ],
               [0.23529412, 0.72361809, 0.47540984, ..., 0.43964232, 0.09042316,
                0.26666667]])
```

```
In [76]: mm_X_test=mm.transform(X_test)
mm_X_test
```

```
Out[76]: array([[0.76470588, 0.52261307, 0.59016393, ..., 0.46497765, 0.16971047,
0.28333333],
[0.23529412, 0.63819095, 0.72131148, ..., 0.51415797, 0.22895323,
0.11666667],
[0.11764706, 0.47236181, 0.62295082, ..., 0.4709389 , 0.25167038,
0.03333333],
...,
[0.          , 0.53266332, 0.57377049, ..., 0.58718331, 0.23207127,
0.01666667],
[0.29411765, 0.62311558, 0.60655738, ..., 0.50670641, 0.06057906,
0.28333333],
[0.17647059, 0.64321608, 0.59016393, ..., 0.4828614 , 0.20712695,
0.1          ]])
```

```
In [77]: mm_lor=LogisticRegression()
mm_lor=mm_lor.fit(mm_X_train,y_train)
```

```
In [78]: mm_y_pred=mm_lor.predict(mm_X_test)
mm_y_pred
```

```
Out[78]: array([0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,
0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [80]: #accuracy
def accuracy(actual,pred):
    return sum(actual==pred)/float(actual.shape[0])
```

```
In [81]: accuracy_score=accuracy(y_test,mm_y_pred)
accuracy_score
```

```
Out[81]: 0.7395833333333334
```

```
In [82]: #precision
print(precision_score(y_test,mm_y_pred))

0.6888888888888889
```

```
In [83]: #recall
print(recall_score(y_test,mm_y_pred))

0.4626865671641791
```

```
In [84]: #auc scores
mm_auc=print(roc_auc_score(y_test,mm_y_pred))
mm_auc

0.6753432835820895
```

## Step-6: [Normalization using StandardScaler and rebuild LOR]

```
In [85]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
ss_X_train=ss.fit_transform(X_train)
ss_X_train
```

```
Out[85]: array([[ -0.85547074,  0.00732864,  0.47259835, ...,  0.88301955,
                -0.65845729, -0.46648591],
                [ 2.46780492, -1.03224482,  0.2585074 , ...,  0.41193433,
                -0.30699915,  1.21865604],
                [ -0.85547074, -0.4022003 , -0.49081095, ...,  0.43739839,
                -0.19585426, -0.8035143 ],
                ...,
                [ -0.85547074, -0.74872478,  0.04441644, ...,  0.76843126,
                -0.78762567, -0.29797171],
                [ 1.56145701,  1.10990656,  0.90078026, ...,  0.28461399,
                2.12917653,  0.71311346],
                [ 0.05087717,  0.73187984, -0.59785643, ..., -0.3265236 ,
                -0.58035548,  0.29182797]])
```

```
In [86]: ss_X_test=ss.transform(X_test)
ss_X_test
```

```
Out[86]: array([[ 2.76992089, -0.5282092 ,  0.15146192, ..., -0.11007904,
                -0.04565848,  0.37608507],
                [ 0.05087717,  0.196342 ,  1.00782574, ...,  0.31007806,
                0.35386232, -0.46648591],
                [ -0.55335477, -0.84323146,  0.36555287, ..., -0.0591509 ,
                0.50706202, -0.8877714 ],
                ...,
                [ -1.15758671, -0.46520475,  0.04441644, ...,  0.93394769,
                0.37488973, -0.97202849],
                [ 0.35299314,  0.10183532,  0.2585074 , ...,  0.24641789,
                -0.78161784,  0.37608507],
                [ -0.2512388 ,  0.22784423,  0.15146192, ...,  0.04270536,
                0.20667045, -0.55074301]])
```

```
In [88]: ss_lor=LogisticRegression()
ss_lor.fit(ss_X_train,y_train)
ss_y_pred=ss_lor.predict(ss_X_test)
```

```
In [90]: #accuracy
def accuracy(actual,pred):
    return sum(actual==pred)/float(actual.shape[0])
```



```
In [91]: ss_accuracy_score=accuracy(y_test,ss_y_pred)
         ss_accuracy_score
```

```
Out[91]: 0.7291666666666666
```

```
In [92]: #precision
         print(precision_score(y_test,ss_y_pred))
```

```
0.6363636363636364
```

```
In [93]: #recall
         print(recall_score(y_test,ss_y_pred))
```

```
0.5223880597014925
```

```
In [95]: #aucscores
         auc_ss=print(roc_auc_score(y_test,ss_y_pred))
         auc_ss
```

```
0.6811940298507462
```


## Step-7. [Plot ROC Curve]

```
In [121]: pred_prob1=mm_lor.predict_proba(mm_X_test)
```

```
In [122]: from sklearn.metrics import roc_curve
         fpr1, tpr1, thresh1 = roc_curve(y_test, pred_prob1[:,1], pos_label=1)
```

```
In [123]: import matplotlib.pyplot as plt
         plt.plot(fpr1,tpr1,linestyle='--',color='orange',label='MinMaxScaler values')
```

```
Out[123]: [ <matplotlib.lines.Line2D at 0x243488be828>]
```



## Step-8. [Comparison with KNN classifier]

```
In [124]: from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=4)
knn=knn.fit(X_train,y_train)
```

```
In [125]: knn_y_pred=knn.predict(X_test)
```

```
In [126]: from sklearn.preprocessing import MinMaxScaler
m=MinMaxScaler()
m_X_train=m.fit_transform(X_train)
m_X_train
```

```
Out[126]: array([[0.05882353, 0.6080402 , 0.63934426, ..., 0.58122206, 0.07884187,
0.11666667],
[0.70588235, 0.44221106, 0.60655738, ..., 0.52608048, 0.13095768,
0.45      ],
[0.05882353, 0.54271357, 0.49180328, ..., 0.5290611 , 0.14743875,
0.05      ],
...,
[0.05882353, 0.48743719, 0.57377049, ..., 0.56780924, 0.0596882 ,
0.15      ],
[0.52941176, 0.7839196 , 0.70491803, ..., 0.51117735, 0.4922049 ,
0.35      ],
[0.23529412, 0.72361809, 0.47540984, ..., 0.43964232, 0.09042316,
0.26666667]])
```

```
In [127]: m_X_test=m.transform(X_test)
m_X_test
```

```
Out[127]: array([[0.76470588, 0.52261307, 0.59016393, ..., 0.46497765, 0.16971047,
0.28333333],
[0.23529412, 0.63819095, 0.72131148, ..., 0.51415797, 0.22895323,
0.11666667],
[0.11764706, 0.47236181, 0.62295082, ..., 0.4709389 , 0.25167038,
0.03333333],
...,
[0.      , 0.53266332, 0.57377049, ..., 0.58718331, 0.23207127,
0.01666667],
[0.29411765, 0.62311558, 0.60655738, ..., 0.50670641, 0.06057906,
0.28333333],
[0.17647059, 0.64321608, 0.59016393, ..., 0.4828614 , 0.20712695,
0.1      ]])
```

```
In [128]: m_knn=KNeighborsClassifier()
m_knn=m_knn.fit(m_X_train,y_train)
```

```
In [129]: m_y_pred=m_knn.predict(m_X_test)
m_y_pred
```

```
Out[129]: array([0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
        0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1,
        1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
        0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1,
        1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0,
        0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1,
        0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [131]: #accuracy
def accuracy(actual,pred):
    return sum(actual==pred)/float(actual.shape[0])
```

```
In [132]: m_accuracy_score=accuracy(y_test,m_y_pred)
m_accuracy_score
```

```
Out[132]: 0.703125
```

```
In [133]: #precision
print(precision_score(y_test,m_y_pred))
```

```
0.5806451612903226
```

```
In [134]: #recall
print(recall_score(y_test,m_y_pred))
```

```
0.5373134328358209
```

```
In [135]: #auc scores
knn_auc=print(roc_auc_score(y_test,m_y_pred))
knn_auc
```


```
0.6646567164179105
```

## Step-9. [Update ROC Curve]

```
In [136]: pred_prob2=m_knn.predict_proba(m_X_test)
```

```
In [137]: from sklearn.metrics import roc_curve
fpr2,tpr2,thresh2=roc_curve(y_test,pred_prob2[:,1],pos_label=1)
```

```
In [141]: import matplotlib.pyplot as plt
plt.plot(fpr1,tpr1,linestyle='--',color='black',label='MinMaxScaler values')
plt.plot(fpr2,tpr2,linestyle='--',color='blue',label='KNN Classifier')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.savefig('ROC',dpi=300)
plt.show()
```



## Step-10. [Regularization]

```
In [139]: from sklearn.linear_model import LogisticRegressionCV
model1=LogisticRegressionCV(Cs=10,cv=4,penalty='l1',solver='liblinear')
model2=LogisticRegressionCV(Cs=10,cv=4,penalty='l2')
```

```
In [142]: model1.fit(mm_X_train,y_train)
model2.fit(mm_X_train,y_train)
```

```
Out[142]: LogisticRegressionCV(Cs=10, class_weight=None, cv=4, dual=False,
    fit_intercept=True, intercept_scaling=1.0, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
    refit=True, scoring=None, solver='lbfgs', tol=0.0001, verbose=0)
```

```
In [143]: rg_y_pred1 = model1.predict(mm_X_test)
rg_y_pred2 = model2.predict(mm_X_test)
```

```
In [144]: #AUC SCORE OF L1
          from sklearn.metrics import roc_auc_score
          l1_auc = roc_auc_score(y_test, rg_y_pred1)
          l1_auc = (' LOR L1 MINMAX AUC', l1_auc)
          l1_auc
```

```
Out[144]: (' LOR L1 MINMAX AUC', 0.6811940298507462)
```

```
In [146]: #AUC SCORE OF L2
          from sklearn.metrics import roc_auc_score
          l2_auc = roc_auc_score(y_test, rg_y_pred2)
          l2_auc = (' LOR L2 MINMAX AUC', l2_auc)
          l2_auc
```

```
Out[146]: (' LOR L2 MINMAX AUC', 0.6851940298507463)
```

## STEP 11 : UPDATE ROC CURVE

```
In [147]: pred_prb7 = model1.predict_proba(mm_X_test)
          pred_prb8 = model2.predict_proba(mm_X_test)
          fpr,tbr,threshold = roc_curve(y_test, pred_prob1[:,1],pos_label=1)
          fpr1,tbr1,threshold1 = roc_curve(y_test, pred_prob2[:,1],pos_label=1)
          fpr2,tbr2,threshold2= roc_curve(y_test, pred_prb7[:,1],pos_label=1)
          fpr3,tbr3,threshold3 = roc_curve(y_test, pred_prb8[:,1],pos_label=1)
```

```

In [151]: plt.plot(fpr, tbr, linestyle='-', color='black', label='LogisticRegression')
plt.plot(fpr1, tbr1, linestyle='-', color='blue', label='KNN')
plt.plot(fpr3, tbr3, linestyle='-', color='red', label='l2')
plt.plot(fpr2, tbr2, linestyle='-', color='orange', label='l1')
plt.annotate(xy=[0.5,0.3],s= auc_ss)
plt.annotate(xy=[0.5,0.2],s= knn_auc)
plt.annotate(xy=[0.5,0.1],s= l1_auc)
plt.annotate(xy=[0.7,0],s= l2_auc)
plt.title('Receiver Operating Characteristic')
plt.legend(loc = 'best')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

